



Embedded RELATED.com



Blogs

- Home
- Blogs
- Forums
- TV
- Jobs
- #IoT
- Tutorials
- Books
- Free PDFs
- Vendors
- Code Snippets
- Login / Register

[Home](#) [Blogs](#) [Forums](#) [TV](#) [Jobs](#) [#IoT](#) [Tutorials](#) [Books](#) [Free PDFs](#) [Vendors](#) [Code Snippets](#)

[Home](#) > [Blogs](#) > [Jason Sachs](#) >

How to Estimate Encoder Velocity Without Making Stupid Mistakes: Part I

Jason Sachs • December 27, 2012 • 30 comments


[Tweet](#)

[Sensors](#) [Measurement](#) [Signal Processing](#)

Here's a common problem: you have a quadrature encoder to measure the angular position of a motor, and you want to know both the position and the velocity. How do you do it?

Sign in

[Sign in](#)

☒ Remember me

[Forgot username or password?](#) | [Create account](#)

You might also like...




Archer Kontrol: 1.27MM Pitch Durable Industrial Connectors

About Jason Sachs

Jason has 25 years of experience in signal conditioning (both analog +

Some people do it poorly -- this article is how not to be one of them.

Well, first we need to get position. Quadrature encoders are **incremental encoders**, meaning they can only measure relative changes in position. They produce a pair of pulse trains, commonly called A and B, that look like this:



This article is available in PDF format for easy printing

You need two of them in order to distinguish between forward and reverse motion. If you need to know absolute position, you can buy a 3-channel encoder, with A and B signals and an index channel that produces a pulse once per revolution in order to signal where a reference position is.

A state machine can be used to examine the A and B pulses and produce up/down counts to a counter of arbitrary width, and the index pulse can be used to initialize the counter to 0 at startup. There are some subtleties here involving noise and failure modes, but let's ignore them, for this article at least.

Most microcontrollers either have built-in encoder peripherals, or interrupt-on-input-change, or both. If you don't have an encoder peripheral, you have to examine the signals yourself and increment/decrement a counter. In any case, let's assume you can get a position count value which is correct based on the encoder signals. Our problem becomes one of taking samples of the position counter and estimating velocity.

Here's a concrete example. Let's consider an object that's rotating but slowing down exponentially due to viscous drag. Here's what the position might look like:



digital) in motion control + medical applications. He likes making things spin.




Popular Posts by Jason Sachs

Understanding and Preventing Overflow (I Had Too Much to Add Last Night)
Chebyshev Approximation and How It Can Help You Save Money, Win Friends, and Influence People
How to Read a Power MOSFET Datasheet
How to Estimate Encoder Velocity Without Making Stupid Mistakes: Part I
Ten Little Algorithms, Part 2: The Single-Pole Low-Pass Filter


Blogs - Hall of Fame




SO YOU WANT TO BE AN EMBEDDED SYSTEMS DEVELOPER
 Steve Branam



INTRODUCTION TO MICROCONTROLLERS
 Mike Silva



IMPORTANT PROGRAMMING CONCEPTS (EVEN ON EMBEDDED SYSTEMS)
 Jason Sachs



HOW FPGA WORK AND WHY YOU'LL BUY ONE
 Yossi Krenin



MSP430 LAUNCHPAD TUTORIAL
 Enrico Garante




ARDUINO ROBOTICS
 Lonnie Honeycutt

Free PDF Downloads



Introduction to Arduino

Advanced Linux Programming



Microcontroller Programming and Interfacing

All FREE PDF Downloads

Here we show the actual continuous position of the object (converted to encoder counts), and the quantized position indicated by the encoder output. Quantized? Yes, both incremental and absolute encoders are digital and have a quantizing effect. A quadrature encoder advertised as "256 lines" or "256 ppr" (pulses per revolution) measures 1024 counts per revolution: quadrature encoders yield 4 counts per cycle, and each cycle corresponds to a "line" (dark mark in an optical encoder) or a pulse measured on an oscilloscope.

In any case, we have lots of choices to obtain velocity from the position count.

The simplest approach is to estimate velocity = $\Delta\text{pos}/\Delta t$: measure the change in position and divide by the change in time. There are two classic options here:

- execute code at fixed time increments (constant Δt), measure position, and take Δpos = the difference between two successive position measurements
- execute code triggered by changes in encoder count (constant Δpos), measure elapsed time, and take Δt = the difference between two successive time measurements

The usual explanation is that the first approach (constant Δt) is better for moderate speeds and the second approach (constant Δpos) is better for low speeds. Texas Instruments has a [user's guide \(SPRUG05A\)](#) discussing this in detail:

General Issues: Estimating velocity from a digital position sensor is a cost-effective strategy in motor control. Two different first order approximations for velocity may be written as:

$$v(k) \approx \frac{x(k) - x(k-1)}{T} = \frac{\Delta X}{T} \quad (1)$$

$$v(k) \approx \frac{X}{t(k) - t(k-1)} = \frac{X}{\Delta T} \quad (2)$$


It goes on to say:

At low speed, Equation 2 provides a more accurate approach. It requires a position sensor that outputs a fixed interval pulse train, such as the aforementioned quadrature encoder. The width of each pulse is defined by motor speed for a given sensor resolution.

[Cue buzzer sound] **WRONG!!!!!!**

The rest of this article will be about understanding why that statement is wrong, why the constant Δpos approach is a poor choice, and (in part II) some more advanced algorithms for estimating position.

First we have to understand what would drive someone to choose the constant Δpos approach. The key to doing so lies in the second sentence I quoted above: "It requires a position sensor that outputs a *fixed interval pulse train*, such as the aforementioned quadrature encoder." If you have a pair of constant frequency square waves in quadrature, and you measure the encoder position in code executing at regular intervals, you might see something like this:



The bottom two waveforms are square waves with a frequency of 1618.03Hz = period of 0.61803 milliseconds (encoder count interval of 1/4 that = 0.154508 milliseconds = 6472.13Hz). The top two waveforms are those same waveforms but sampled at 0.1 millisecond intervals.

You'll note that the time between encoder state changes perceived by those sampled waveforms varies irregularly between 0.1 millisecond and 0.2 millisecond -- as it should be; we can't possibly measure 0.154508 millisecond intervals accurately, with code that executes at 0.1 millisecond granularity. Code that uses the constant Δt approach to measure velocity, with $\Delta t = 0.1$ milliseconds, would measure a velocity of either 0 or 10000Hz for this waveform at each sampling interval. **The velocity resolution, for code that uses calculates $\Delta\text{pos}/\Delta t$ using constant Δt , is $1/\Delta t$.**

This leads to three usual impulses:

1. Why don't I just use a larger Δt ? Then my velocity resolution is smaller.
2. Why don't I just use a low-pass filter on the $\Delta\text{pos}/\Delta t$ calculations?
3. Why don't I just use the constant Δpos approach?

We'll discuss the first two of these later. The constant Δpos approach (i.e. get interrupts at each encoder state change and measure the time), *for fixed frequency waveforms*, is actually the optimal approach. If I have a timer with 0.1 microsecond accuracy, and I get an interrupt on each encoder state change, then my resolution for velocity is not constant, like the constant Δt approach, but instead depends on the frequency. For this particular set of waveforms representing an encoder state frequency of 0.154508 milliseconds per state change, a 0.1 microsecond timer resolution would either measure 0.1545 or 0.1546 milliseconds between state changes = 6472.5Hz or 6468.3Hz -- for a velocity resolution of about 4.2Hz. The slower our waveforms, the more precise our velocity resolution can be

using the constant Δpos approach.


What's wrong with that?

Well, for one thing, as the waveforms get faster, the precision gets worse, and the more interrupts we would get, leading to increased CPU usage. There's also the case of vibration around zero speed: you then have to distinguish between positive and negative position changes, and the velocity estimates in that case may be excessive. Those problems are pretty obvious.

The more subtle issue here is that the optimality of the constant Δpos approach is true *only* for fixed frequency waveforms. If the frequency varies with time in an unknown manner, the constant Δpos approach only gives us the average frequency between instants at which the encoder state changes, which may not be what we want. I also said "the slower our waveforms, the more precise our velocity resolution can be" -- and *precise* is not the same as *accurate*.

Let's go back to our earlier example of exponential decay.

Here's what an estimated velocity waveform looks like when using the constant Δpos approach:




The blue is the actual velocity during this time, and the green is the estimated velocity. Notice how the green waveform is always above the blue waveform. In general this is always true, due to the backward-looking property of any real algorithm: when velocity is decreasing, any estimate of velocity will lag behind and be too high, and when velocity is increasing, any estimate of velocity will lag behind and be too low.

Now remember we only update the estimate when we get a change in encoder position. This is a problem with the constant Δpos approach. Suppose an object we're measuring is moving along at a certain slow but constant positive velocity: the encoder count updates periodically, and we'll calculate this velocity correctly. Now the object stops. No more encoder counts. But until we get another count, we won't update the velocity estimate! We never estimate a velocity of zero! We can "fix" this problem with a hack: set a timeout and when the timeout expires, estimate a velocity of zero. These sorts of fixes in signal processing are generally very bad: they're hard to analyze and you can't predict when the fix is going to backfire.

Aside from the zero velocity problem, the issue in general is due to the area between the two waveforms: actual and estimated velocity. In a good velocity estimator, this area will always be bounded by a number that's proportional to the acceleration rate. Here the area is larger the closer we get to zero velocity. Why do we care? Because the area is the integral of velocity, which is position: if you integrate the velocity estimate over time, you should get back the same value of position as you started with. **The constant Δpos approach does not produce bounded position drift when you integrate the velocity estimate to obtain position.** Does this matter in all systems? No, but it has consequences whenever you're using integrals of estimated velocity, for example in a PI controller.

Another closely related issue is the time at which velocity estimates are obtained: they aren't at constant time intervals, but rather occur at each change in position count. Aside from the inconvenience of getting lots of interrupts at unpredictable times with high velocities, this means that they're not directly compatible with regularly-sampled control systems. There's a lot of valuable theory behind a control system that executes at regular intervals, and there are few reasons to deviate from this approach. If you have a PI controller that runs at regular intervals, and it gets velocity estimates from the constant Δpos approach, there will always be a time delay between the velocity estimate update, and the point in time at which it is used: the velocity estimate is effectively resampled at the rate at which it's used. This time delay is not constant and causes additional area between the real and estimated velocity waveforms:



It's not much but it causes additional error. (Look closely at this graph and the previous one.)

All of this discussion so far has assumed one thing about the encoder: that it's a perfect measuring device. We've quantized the real position into exactly equal intervals. A real position encoder has errors due to manufacturing tolerances.

The main mistake of using the constant Δpos approach is assuming an encoder is perfect. It's not. Let's look at an encoder I've used before, the Avago HEDS-55xx/56xx series. If you look at p. 4-6 of the datasheet, you'll see specifications, diagrams, and definitions for the errors in the encoder waveforms. Here are the position accuracy specifications:

Encoding Characteristics

Part No.	Description	Sym.	Min	Typ.*	Max.	Units
HEDS-5500	Pulse Width Error	ΔP		7	45	°e
HEDS-5600	Logic State Width Error	ΔS		5	45	°e
(Two Channel)	Phase Error	$\Delta \Phi$		2	20	°e
	Position Error	$\Delta \Theta$		10	40	min. of arc
	Cycle Error	ΔC		3	5.5	°e
HEDM-5500	Pulse Width Error	ΔP		10	45	°e
HEDM-5600	Logic State Width Error	ΔS		10	45	°e
(Two Channel)	Phase Error	$\Delta \Phi$		2	15	°e
	Position Error	$\Delta \Theta$		10	40	min. of arc
	Cycle Error	ΔC		3	7.5	°e
HEDS-5540	Pulse Width Error	ΔP		5	45	°e
HEDS-5640	Logic State Width Error	ΔS		5	35	°e
(Three Channel)	Phase Error	$\Delta \Phi$		2	15	°e
	Position Error	$\Delta \Theta$		10	40	min. of arc
	Cycle Error	ΔC		3	5.5	°e
	Index Pulse Width	P_o	55	90	125	°e
	CH.I rise after -40°C to +100°C CH.A or CH. B fall	t_1	-300	100	250	ns
	CH.I fall after -40°C to +100°C CH.A or CH. B rise	t_2	70	150	1000	ns
HEDM-5540	Pulse Width Error	ΔP		10	45	°e
(Three Channel)	Logic State Width Error	ΔS		10	45	°e
	Phase Error	$\Delta \Phi$		2	15	°e
	Position Error	$\Delta \Theta$		10	40	min. of arc
	Cycle Error	ΔC		6	12	°e
	Index Pulse Width	P_o	50	90	130	°e
	CH.I rise after -40°C to +85°C CH.A or CH. B fall	t_1	200	1000	1500	ns
	CH.I fall after -40°C to +85°C CH.A or CH. B rise	t_2	0	300	1500	ns

Note: See Mechanical Characteristics for mounting tolerances. *Typical values specified at $V_{CC} = 5.0\text{ V}$ and 25°C .

The encoder with the best specifications here is the HEDS-5540/5640 encoder, which specifies state width error $\Delta S = 35$ degrees electrical maximum, 5 degrees typical. What does this mean for us if we're using the Δpos approach to estimate velocity?

Each encoder state is nominally 90deg electrical, so a ± 35 degrees electrical error is a $\pm 39\%$ error; ± 5 degrees electrical error is $\pm 5.5\%$ error. We calculate $v = \Delta\theta/\Delta t$. The time measurement Δt is determined by the accuracy of the CPU clock, which for a crystal oscillator is typically 50ppm or less – essentially negligible. So the error in velocity is directly proportional to the error in $\Delta\theta$ – typically 5.5% and up to 39% error from one encoder state to the next!

We could use the time between each edge and the 4th-previous edge instead; this is the cycle time C specified in the datasheet. The error ΔC is 3 degrees typical, 5 degrees maximum – with a 360 degree total cycle this works out to 0.8% typical, 1.4% maximum.

Note that these errors aren't DC errors, they appear as jitter from one state or one cycle to the next. So if we average these velocity estimates then the error is attenuated, which is fine for high speeds, but ruins the advantage of using the time between one encoder edge to the next – but this is the whole point of using the constant Δpos approach!

So that's why I don't recommend using it. If you're using an encoder, and you can analyze the datasheet for it to determine the typical or worst-case errors, and you're fine with these numbers, then go ahead and use constant Δpos if you must.

Well, if constant Δpos is out, what's left?

Let's be fair and look at the constant Δt approach with the same kind of analysis. Here we must make an estimate. Let's assume the position count we measure represents the center of the encoder state. With a perfect encoder, sampling the position at any given instant can

be off by up to 45 degrees electrical: the actual position could be at either edge of the pulse. This means that between two successive measurements we could be off by 90 degrees electrical.

Examples for which this is true: we make one measurement at the beginning of one encoder state, and the next measurement near the other end of that same encoder state, so the real change in position is 90 degrees electrical, and the estimated change in position is 0. Or we make one measurement just before the transition of an encoder state, and the other measurement just after the following transition, so the real change in position is 90 degrees and the estimated change in position is 2 counts = 180 degrees electrical. With an imperfect encoder that has a worst-case state width error of 35 degrees electrical, the worst-case value for real change in position in these two cases are $90+35 = 125$ degrees electrical for an estimated change in position of 0, or $90-35 = 55$ degrees electrical for an estimated change in position of 180 degrees; both cases add that worst-case error of the imperfect encoder, for a total worst-case error in Δpos of 125 degrees.

This error is a property of the *encoder*, not of the algorithm used.

The point I want to make here is that because encoders are imperfect, there will be a position error that we just have to expect and accept. With the constant Δt approach, this translates into a constant worst-case speed error, independent of speed. We have a choice of which Δt we want to use, which gets us back to two points I raised earlier:

1. Why don't I just use a larger Δt ? Then my velocity resolution is smaller.
2. Why don't I just use a low-pass filter on the $\Delta\text{pos}/\Delta t$ calculations?

The choice of Δt leads to a tradeoff. If we choose a large Δt , we get very low speed error, but we will also get a low bandwidth and large time lag. If we choose a small Δt , we get very high bandwidth and small time lag, but we will also have high speed estimation error.

If I were using the basic approach for estimating angular velocity, I would take the position estimates made at the primary control loop rate, take differences between successive samples, divide by Δt , and low-pass-filter the result with a time constant τ that reduces the bandwidth and error to the minimum acceptable possible. (A 1-pole or 2-pole IIR filter should suffice; if that kind of computation is expensive but memory is cheap, use a circular queue to store position of the last N samples and extend the effective Δt by a factor of N .)

If you're using the velocity estimate for measurement/display purposes, then you're all done. If you're using the velocity estimate for control purposes, be careful. It's better to keep filtering to a minimum and accept the velocity error as noise that will be rejected by the overall control loop performance.

Is the best we can do? Not quite. We'll discuss some more advanced techniques for velocity estimation in Part II of this article.

You might also like... (promoted content)




Archer Kontrol: 1.27MM Pitch Durable Industrial Connectors

Connectors

Comments

Comments

 Write a Comment Select to add a comment

[-]

Comment by **kevbo** • February 15, 2018 3

I'll add my 2 cents: Not usually quadrature encoded, but similar issues arise with pulse-output flow meters: Flow rate is pulses/deltaT, and that is what is often needed for process control.

I have seen three different programmers** that calculated the flow rate that way, then went on to integrate those flow readings to obtain a volume. (which is often needed for maintenance intervals or auditing, etc) Uhhhh....

In case you are not seeing the issue: pulses/deltaT is a derivative,* which they would then integrate. Both of which have potential error sources. The correct way to totalize the volume is just to count the pulses, because each pulse represents a fixed volume.

Finally, another big error I have seen is using a floating point variable as a volume accumulator. This works great until you accumulate enough volume that at low flow rates your per-interval volume is less than the least significant digits of your mantissa...which causes you to increment the volume by zero. This never gets caught in commissioning tests, because they never let it run long enough to push the accumulator exponent high enough to matter. Double precision floats help, or regular transfer of accumulated sub-total to a "master accumulator" and reset of the sub accumulator are band-aids for those that insist on floating point. I prefer to use integers. (several if needed). When the raw counts are divided to convert to units, I save the remainder (modulo) and add it to the next interval's count value, thus avoiding truncation error. Eventually the added remainders add up to a whole unit (gallons say) and the error goes back to zero.

*OK, it is not strictly a continuous time derivative, but same thing for this discussion.

**Actually one of the programmers was using the pulse input to an electronic speedometer, converting to speed, then integrating for distance. It is the same issue, just speed and distance, rather than flow and volume.

Reply Reply 

[-]

Comment by **ilikecheese** • November 8, 2013 2

You can't leave us hanging like this!! Where is part 2?

Reply Reply 

[-]

Comment by **jms_nh** • June 25, 2014 3

Sorry I missed noticing your comment -- I did post part 2:

<http://www.embeddedrelated.com/showarticle/530.php>

Reply Reply 

[-]

Comment by **Evan B** • January 7, 2013

very interested to read part 2.

1

Reply

Reply

[-]

Comment by **stephaneb** • June 9, 2016

First image seems to be gone.

1

Reply

Reply

[-]

Comment by **jms_nh** • June 9, 2016

thanks, the url changed >:(-- i'll fix

0

Reply

Reply

[-]

Comment by **Kenny_L** • July 6, 2018

Really fantastic article here. Thanks very much for the time to explain and write all this. You taught us a lot. I'll certainly be using these techniques shown after seeing and understanding the presented issues.

1

A quick note for now for the part that says "a 0.1 microsecond timer resolution would either measure 0.1545 or 0.1546 milliseconds between state change". That should be 100 microsecond timer resolution, right? It's just a minor typo maybe. Thanks again! Genuinely appreciated for this article, and for part 2 as well.

Reply

Reply

[-]

Comment by **WestfW** • January 3, 2013

I have to ask - why are we trying to re-derive the velocity from the position data (introducing extra layers of error), when we could simply compute it in the same pin-change ISR that's calculating position in the first place. Not counting "stopped", the velocity is just $1/dt$ where dt is the time delta between ANY two signal transitions. Yes, this would be subject to the same errors, but not more than a position-based velocity, and it seems like it would be easier to "smooth" the velocity calculations at edge times rather than sample times of the position.
(This is probably not applicable if you have an encoder peripheral...)

0

Reply

Reply

[-]

Comment by **jms_nh** • January 4, 2013



WestfW: The method you suggest *is* the constant delta-pos method. You don't actually compute a delta-pos, it's a constant: +/-1 encoder state. (Remember you have to keep track of direction.)

👍 0

With regard to smoothing: filtering really doesn't work well when the rate of execution is variable and unpredictable.

Reply ↩

Reply ↩

[-]

Comment by [sparkygsx](#) • January 3, 2013



What I've done in the past, is recording both the current position (counting up or down as appropriate), and the timestamp of the last step forward and reverse. The code that runs the control loop, at a fixed and relatively slow interval, reads those 3 parameters, and keeps a copy of the recorded values of the previous cycle. Basically, you take the time between the most recent event and the last event of the previous cycle, and divide by the number of events, to determine the average time between events since the last update of the control loop. If the recorded previous values are only overwritten when there are new events, you've covered the case where the input frequency is lower than the loop update frequency, and the time since the last recorded event gives you an upper bound on the speed, so that it will asymptotically go to 0 when no events are received.

👍 0

Reply ↩

Reply ↩

[-]

Comment by [sparkygsx](#) • January 4, 2013



What I forgot to say: the method I described sort of seamlessly transitions between the delta-time and delta-position modes, while minimizing the main source of noise in both; you will never measure the time of a number of events +/- 1 which causes problems at low speed, or intervals of events of +/- 1 cycle, which causes problems at high speed.

👍 0

Basically, it's the third obvious way: it calculates delta position / delta time, where neither is fixed, but both are maximized.

The timestamp can be very high resolution, because it can be a free running counter at the maximum clock speed.

It is, of course, still important to minimize the interrupt latency, because this limits the input frequency, and especially the latency jitter, because this determines the accuracy of the timestamps.

Reply ↩

Reply ↩

[-]

Comment by [jms_nh](#) • January 4, 2013



I agree that adding new information of edge times between the regularly-sampled position measurements is a potential improvement. But that information needs to be incorporated carefully. The method you described really needs to be analyzed from a signal processing perspective or simulated or compared with more traditional methods to show it leads to improved accuracy and does not introduce error. (especially considering the nasty cases of vibration around an encoder edge, or sudden acceleration/deceleration)

👍 0

In part II I'll discuss some more sophisticated methods which also have some solid theoretical background.

Reply ↩

Reply ↩

[-]

Comment by [sparkygsx](#) • January 4, 2013



0

Where did my 2 comments go? I wrote a fairly long piece about a method I've used before to avoid the problems mentioned, by counting pulses and recording timestamps of the last pulse of each control loop cycle, which would be a delta position / delta time method with neither delta as constant, including some simple fixes for the zero speed and very low speed cases. I had put in quite a bit of effort, and it seems like you just removed it!

Reply

Reply

[-]

Comment by [jms_nh](#) • January 4, 2013

0

They weren't removed! I think the forum software automatically holds multiple comments for approval. I just approved them so there they are.

Reply

Reply

[-]

Comment by [sparkygsx](#) • January 4, 2013

0

Weird, I could see the messages for about half an hour before they disappeared. No worries though.

I do agree that the method I described would have to be analysed. I've used it before for wideband frequency measurement, when the back-and-forth stepping problem does not exist. In that case, it solves the problems you mentioned in both modes with either fixed delta-pos or fixed delta-time, as well as potential cross-over issues if you would implement both. I couldn't find any potential new issues, besides possibly a slight increase in CPU load.

I suggested recording the timestamps for the last up- and downcounting events separately and using either one depending on the net direction of rotation, and using the delta-position as opposed to the number of events to suppress speed overestimation caused by back-and forth stepping.

Of course, some problems, like errors in the phase shift of the channels and such would still cause inaccuracies. One possible improvement could be to measure the time between two occurrences of the same edge event; this would at least eliminate sensor placement errors. Correcting for imperfect wheel slot placement will probably be a bit more difficult. If you have an index signal, you could store the angle of the separate events after some kind of calibration procedure, but I'd think that would be fairly complicated and CPU- and memory intensive.

I'm rather curious about your solutions, though.

Reply

Reply

[-]

Comment by [Søren Friis](#) • January 14, 2013

0

Very interesting indeed. Looking forward to part II :) !

Reply

Reply

[-]

Comment by [p1r](#) • January 22, 2013

0

Interesting reading definitely want to read part two. When driving a motor: instead of using quadrature encoder you could also use single channel encoder and direction is indirectly known since you know the direction you are driving the motor. I have seen this in consumer electronics where saving one hall sensor can make sense.

Reply

Reply

[-]

Comment by **sparkygsx** • January 23, 2013

That may work in some cases, but there are a few things to consider: first, in some cases the motor could rotate in the opposite direction if the load can supply a torque to turn the motor, and second, if the motor is stopped, a small vibration could cause the sensor to generate pulses if there is insufficient hysteresis.

2

Reply

Reply

[-]

Comment by **Tim Wescott** • November 19, 2013

I have seen some really bad things happen by trying to rely on this method. It can work quite well if you can be absolutely sure that the motor is only rotating in one direction, but I have seen it tried in a control loop that needed to work at low velocities and through speed reversals of the motor, and it was an absolute disaster.

0

Reply

Reply

[-]

Comment by **jms_nh** • January 25, 2013

^^ what Sparky said. :)

0

Reply

Reply

[-]

Comment by **mrajesh_9** • February 24, 2013

I am using TMS320F2812 for sensing position of switched reluctance motor. My motor has four hall sensor each displaced by 7.5 degree. and each sensor gives only six pulse per revolution. signals s1, s2, s3 and s4 are displaced by 7.5deg with 30deg on period and 30deg off period. s1-s3 and s2-s4 forms pair of quadrature signals. I dont know which method I should use for sensing position. I had used simulink interface but simulink trigger is not working with kit. I tried capture and timer .also

0

Reply

Reply

[-]

Comment by **effgee** • February 24, 2013

what's the outlook for Part 2 ?
Taking it just this far is a bit of a tease !

0

Reply

Reply

[-]

Comment by **volkhen** • March 2, 2013

Your blog is something that I was searching for a long time. I really appreciate that someone's so experienced is sharing his knowledge. I'm looking forward to Part II. In the mean time I'll change my PID algorithm according to your advices in my dc motor servo drive.

0

Reply

Reply

[-]

Comment by **mikecook** • April 1, 2013

Excellent article. As usual it was very informative. Maybe I can use it to explain to my bosses that things are harder than they appear when it comes to useful velocity measurements. I look forward to part 2.

0

Reply

Reply

[-]

Comment by **SANEESH MP** • August 19, 2013

Could you share more efficient method to estimate low speed.

Thanks
saneesh

0

Reply

Reply

[-]

Comment by **Max25** • April 30, 2016

I can't understand why you wrote about problems to measure delta t, CPU load, interrupts, timer with 0.1 microsecond accuracy and etc? All motorcontrol microcontrollers can measure time between encoder edges by hardware, without program code. You need to get ready result from register of latched time when you need to calculate your speed PID.

0

Reply

Reply

[-]

Comment by **jms_nh** • April 30, 2016

All motorcontrol microcontrollers can measure time between encoder edges by hardware, without program code.

And that's not a great way to estimate velocity. It works OK for slow rotation, bad for fast rotation, and it is sensitive to the manufacturing tolerances of the encoder edges, which I mentioned in this article and another one (<https://www.embeddedrelated.com/showarticle/444.php>)

0

Reply

Reply

[-]

Comment by **Markus47** • September 28, 2016

Will there be a part III?

0

Reply

Reply

[-]

Comment by **jms_nh** • October 8, 2016

Part III requires some more investigation that unfortunately I don't have time to do in at least the next 12-18 months. Sorry.

0

Reply

Reply

[-]

Comment by **Bandicoot** • October 11, 2016

👍 2

No worries. I'm looking forward until you find the the time.

I'm really curious how you close the gap between digital encoders, PLLs and observers. PLLs seem to be related to angle tracking observers which use resolvers but not digital encoders, as far as I understand it.

Any recommendations on papers on digital encoders in combination with observers or anything else on the topic in the mean time?

Reply ↩

Reply ↩

To post reply to a comment, click on the 'reply' button attached to each comment. To post a new comment (not a reply to a comment) check out the 'Write a Comment' tab at the top of the comments.

Quick Links

Home
Blogs
Forums
TV
Jobs
#IoT
Tutorials
Books
Free PDFs
Vendors
Code Snippets
comp.arch.embedded

About EmbeddedRelated.com

Advertise
Contact
Privacy Policy
Terms of Service
Cookies Policy

Social Networks



The Related Media Group

